

## FTEC 4003 Project Report

Name: IP Shing On

SID: 1155109011

Name: WONG Shing

SID: 1155109027

### Task 1:

#### Background:

We use Visual Studio with anaconda installed to do task 1. We import the corresponding libraries to evaluate each of the classification methods. We choose Bagging as the sample of Ensemble Method.

#### Experimental evaluations:

Performance of all the methods:

Decision Tree	0.3057297756989476
k-Nearest Neighbor	0.18764056977615937
SVM	0.07607090103397342
Bayes	<b>0.42970297029702964</b>
Bagging	<b>0.4298139381261697</b>

Therefore, the best 2 methods are Bayes and Bagging.

#### Methods:

##### Decision Tree:

Settings:

```
DecisionTreeClassifier(criterion= 'gini', splitter= 'best', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, max_features= "sqrt",  
random_state=50)
```

For random\_state:

max_depth	random_state
None	random_state(0) = 0.28996442815565376 random_state(1) = 0.290893489220208 random_state(2) = 0.2869397383658589 random_state(3) = 0.30073899539466636 random_state(4) = 0.28858191236779623 random_state(5) = 0.2835724994645534 random_state(6) = 0.28990472112193555

	random_state(7) = 0.29394626092586595 random_state(8) = 0.2931498010538768 random_state(9) = 0.291363830696879 random_state(10) = 0.2973580062038721 <b>random_state(50) = 0.3057297756989476</b> random_state(100) = 0.3007646559048428 random_state(150) = 0.2835757057313944 random_state(200) = 0.2835757057313944
--	---

We find that random\_state=50 gives the highest performance.

Finding the best max\_depth: (random\_state=50)

max_depth	performance
None	<b>0.3057297756989476</b>
< 8	0
8	0.0008710801393728222
9	0.0021743857360295715
10	0.005183585313174947
20	0.12378184305009403
50	0.29053177691309984
55	0.3000319863524896
56	0.30554964916011057
57	<b>0.3057297756989476</b>
58	0.3057297756989476
70	0.3057297756989476

We first find that the performance of max\_depth<20 is very low, so we jump to test max\_depth=50. Then, we find that when max\_depth=57, the performance is the same as max\_depth=None. We further increase the max\_depth but the performance seems to be maximized. It means that the fully developed tree has a max\_depth of 57. With random\_state=50, we have a performance of 0.3057297756989476. That's the best performance we can get after tuning the parameters.

## k-Nearest Neighbor:

Settings:

```
KNeighborsClassifier(n_neighbors=1, weights= 'distance')
```

We set weights to be 'distance' instead of 'uniform' because we put the weights into consideration, which can provide a better performance.

To find the n\_neighbors:

n_neighbors	Performance	
	weights= 'distance'	weights= 'uniform'
1	<b>0.18764056977615937</b>	<b>0.18764056977615937</b>
2	0.18475550254797787	0.06652126499454744
3	0.14000562271577172	0.13037166085946572
4	0.13406688866367455	0.045358447049778976
5	0.1042622950819672	0.09061488673139159
6	0.09632488672596073	0.030923694779116464
7	0.07237549265496238	0.05272013460459899

We can see that the performance decreases while n\_neighbors increases. Finally, we set n\_neighbors=1 and weights= 'distance'. That means there are a total of 2 elements in each group.

## SVM

Settings:

```
SVC(C=5, kernel='linear', random_state=0)
```

We choose 'linear' to be the kernel because the dataset can be linearly classified. Linear classification, which is the main topic of this course, can do classification fast and accurately.

Since running a large amount of data with SVM requires plenty of time, we sample the data from 300000th.

C	Performance(linear)	Performance(rbf)
1	0.07480096278466951	0.0021715526601520083
2	0.07507396449704141	0.010559662090813094

3	0.0747871158830063	0.013033424427159974
4	0.07417022065640647	0.013445378151260505
5	<b>0.07607090103397342</b>	0.014276716355238295
6	0.07601476014760147	0.015107007973143096
7	0.07438934122871947	0.015107007973143096
8	0.07408779403593257	0.015107007973143096
9	0.07447202667654687	0.01510383889238515
10	0.07548566142460685	0.01510383889238515
20	0.07445823300611225	0.015100671140939596
50	0.07538802660753881	0.015100671140939596

Therefore, we find that c=5 with a linear kernel gives the best performance.

### Bayes:

Settings:

CategoricalNB(alpha=3, fit\_prior=False, class\_prior=None)

There are five models of Naive Bayes including Gaussian Naive Bayes, Multinomial Naive Bayes, Complement Naive Bayes, Bernoulli Naive Bayes and Categorical Naive Bayes.

Referring to the formulas in the lecture notes  $\Pr[A_i = X_i | Y = Yes] = \frac{N_{i, Yes} + 1}{N_{Yes} + c}$ , we first eliminate Gaussian Naive Bayes and Complement Naive Bayes. Then, we eliminate Bernoulli Naive Bayes because the features are not necessarily binary-valued. For instance, age, region\_code, Annual\_Premium consists of a range of values instead of only binary-value. For Multinomial Naive Bayes, since it is mainly used for text classification where data are represented as word vector counts, it is eliminated. Finally, looking at Categorical Naive Bayes, it is good for classification with categorically distributed discrete features. In our train dataset, each feature like Driving\_License with 0 and 1 as well as Vehicle\_Damage with 'No, and 'Yes' has its own categorical distribution. Therefore, we finally choose to implement Categorical

Naive Bayes with the formula  $P(x_i = t | y = c; \alpha) = \frac{N_{tic} + \alpha}{N_c + \alpha n_i}$ , where 'alpha' (Additive Laplace smoothing parameter (0 for no smoothing)) is behaving like 'c' in the formula of the lecture note. When alpha becomes zero, the whole expression becomes zero when  $N_{tic}$  is zero.

For parameter `fit_prior`, we just keep the classes with a uniform prior, so we set it to `False`. For parameter `class_prior`, since the classes have uniform prior probabilities, we set the `class_prior` to `None`.

Finding a suitable alpha:

alpha	Performance
0 (without smoothing/without Laplace)	0.38751625487646296
1	0.42792041791708346
2	0.42943904593639576
<b>3</b>	<b>0.42970297029702964</b>
4	0.42902588854760865
5	0.4294391244870041
6	0.42945321461735947
7	0.429226235948925
8	0.4292177705096757
9	0.42910081743869205
10	0.4291784702549575
15	0.42826878157423004
20	0.4284939236111111
30	0.42786716506853034
40	0.42776694823478456
60	0.4278213802435724
100	0.4276101584545257
200	0.4278460278460279
500	0.4221806429664909

The performance without Laplace is not good, with a performance of 0.3735070575461455. Once the Laplace is set, that is setting a value to alpha, the performance improves. Overall, the nearer the alpha to be 3, the better the performance of the Naive Bayes Algorithm. So, we finally set alpha to be 3.

To explain the alpha, we may first look at the two formulas:  $\Pr[A_i = X_i | Y = Yes] = \frac{N_{iYes} + 1}{N_{Yes} + c}$

-(1) and  $P(x_i = t | y = c; \alpha) = \frac{N_{tic} + \alpha}{N_c + \alpha n_i}$  -(2). The former formula that is in the lecture notes has numerator+1 and denominator+c. The latter formula that is in the implementation has numerator+alpha and denominator+ni\*alpha.

In the training dataset, we have only 2 classes to predict, so c=2.

ni is the number of available categories of feature i. For instance, in the training dataset, previously\_insured has 2 categories, so ni for previously\_insured is 2.

We have found two relations:

1. (1) will be greater than (2) when the raw probability increases, for any alpha with a fixed value on ni.
2. For a fixed raw probability, an increase of ni will decrease (2).

1. Let alpha = 3, ni=2, c=2

Ni,Yes	NYes	lecture result	Implementation result
1	10	0.167	0.25
2	10	0.25	0.3125
3	10	0.333	0.375
4	10	0.4167	0.4375
5	10	0.5	0.5
6	10	0.5833	0.5625
7	10	0.667	0.625

2. Let alpha = 3, c=2

Ni,Yes	NYes	ni	lecture result	Implementation result
1	10	2	0.167	0.25
1	10	3	0.167	0.211
1	10	4	0.167	0.182
1	10	5	0.167	0.16

For the training dataset, the attributes are having different numbers of categories, resulting in different numbers of ni. Alpha is tuned to balance the differentiation brought by ni and the raw probability. Alpha is tuned to make a better decision. For instance, alpha is tuned to make (2) equals to (1).

## Baggings:

Settings:

```
bayes = CategoricalNB(alpha=3, fit_prior=False, class_prior=None)
```

```
BaggingClassifier(base_estimator=bayes, n_estimators=15, random_state=4,  
bootstrap=True)
```

We have used DecisionTree, KNN, Bayes and SVM as the base\_estimator.  
We find that Bayes performs the best.

Random state is set as to eliminate the randomness of the performance. And  
we find that the performance slight improves when random\_state is set to 4.

bootstrap=True means that samples are drawn with replacement.

Finding the suitable n\_estimators (number of classifiers):

n_estimators = 1	0.42861072902338376
n_estimators = 2	0.4289800858180218
n_estimators = 3	0.4295309403215151
n_estimators = 4	0.42941273597886515
n_estimators = 5	0.42910612065169523
n_estimators = 6	0.4292081659604909
n_estimators = 7	0.429137323943662
n_estimators = 8	0.4291137151345107
n_estimators = 9	0.4291293108189868
n_estimators = 10	0.42945932566965506
n_estimators = 11	0.42950657351889543
n_estimators = 12	0.4297502475519858
n_estimators = 13	0.42962473863761413
n_estimators = 14	0.42962473863761413
n_estimators = 15	<b>0.4298139381261697</b>
n_estimators = 16	0.4295619634602685
n_estimators = 17	0.42962473863761413
n_estimators = 18	0.42952281358357647
n_estimators = 19	0.42949105914718017

n_estimators = 20	0.4295937465595068
n_estimators = 21	0.4296011004126547
n_estimators = 22	0.42951469131726644
n_estimators = 23	0.4292473591549295
n_estimators = 24	0.4292629262926293
n_estimators = 25	0.4292629262926293
n_estimators = 26	0.4293493207194324
n_estimators = 27	0.4293020953638014
n_estimators = 28	0.4293493207194324
n_estimators = 29	0.4293020953638014
n_estimators = 30	0.4293493207194324
n_estimators = 40	0.4295693306198779
n_estimators = 60	0.42948294829482947
n_estimators = 100	0.42946742957746475
n_estimators = 200	0.42949105914718017
n_estimators = 500	0.42959295929592967

Though the performance keeps fluctuating, we find the best performance at n\_estimators=15. Unlike Bayes, there isn't a clear tendency of the performance while the parameter is increasing. n\_estimators=15 means that original data is split into 15 subsets and each of them is classified under Naive Bayes. n\_estimators=15 is concluded solely based on the experimental result.

### **Comparison of the two best methods:**

Bagging has a slightly better performance than Bayes, with 0.4298139381261697 > 0.42970297029702964. Theoretically, this is because bagging repeats the classification a number of times and the probability of having a wrong decision is highly decreased. Mathematically, we can refer to the lecture notes.



- Suppose there are 25 base classifiers
  - Each classifier has error rate  $\epsilon = 0.35$
  - Assume errors made by classifiers are uncorrelated

- Probability that the ensemble classifier makes a wrong prediction

$$\Pr[X \geq 13] = \sum_{i=13}^{25} \binom{25}{i} \cdot \epsilon^i (1 - \epsilon)^{25-i} = 0.06$$

Among the methods, Bagging with Bayes as the base estimator gives the best performance. This is because Bayes itself is already performed well. Under Bagging, the performance will be further better off when applying the Bayes classification again and again .

## Task 2:

### Background:

We use Visual Studio with anaconda installed to do task 2. We import the corresponding libraries to evaluate each of the classification methods.

### Experimental evaluations:

Performance of all the methods:

Decision Tree	0.6111111111111111
k-Nearest Neighbor	0.1744421906693712
Bayes	0.5109374999999999
SVM	0.1602671118530885
Bagging	<b>0.630071599045346</b>

Therefore, the best method is Bagging with Decision Tree.

### Methods:

#### Decision Tree:

Settings:

`DecisionTreeClassifier(min_samples_split=257, random_state=0)`

To tune the parameters, we first use GridSearchCV to find the approximate range of the best parameters. Then, we tune the parameters by ourselves.

First, set random state = 0 to eliminate the randomness of the performance.

Finding the min\_samples\_split:

GridSearchCV range	Result (min_samples_split)	Performance
range(2,100,1)	99	0.5846867749419954
range(2,200,1)	<b>199</b>	<b>0.5875</b>
range(2,300,1)	<b>299</b>	<b>0.5898942420681551</b>
range(2,400,1)	398	0.5707317073170731
range(2,500,1)	449	0.5707317073170731

We then search with a shorter range. Starting from the second group to the third group.

GridSearchCV range	Result (min_samples_split)	Performance
range(190,220,1)	218	0.5873417721518988
range(221,251,1)	221	0.5873417721518988
range(252,282,1)	<b>281</b>	<b>0.5898942420681551</b>
range(283,300,1)	299	0.5898942420681551

We can see that the third group gives the best performance. We then further fine tune min\_samples\_split by ourselves.

min_samples_split	Performance
252 - 255	0.5880861850443599
256	0.5880861850443599
<b>257</b>	<b>0.6111111111111111</b>
258 - 268	0.6111111111111111
269 - 282	0.5898942420681551

We find the best performance at min\_samples\_split=257, with performance of 0.611.

We then modify the random\_state. However, we find that random\_state does not affect the performance. We try random\_state with values from 1 to 17, but the performance is still 0.611. On the other hand, for max\_depth, since we do not set the max\_depth, it will be None by default which means the tree is expected to be fully developed. We further test this by putting values to max\_depth.

With max\_split=257, random\_state=0

max_depth	performance
1	0
2	0.5077844311377245
3	0.4131455399061033
4	0.5361366622864652
5	0.5944645006016848
6	0.5944645006016848
7	0.6060606060606062
8	0.6095017381228274
9	0.6095017381228274
10	0.6095017381228274
11	0.6095017381228274
<b>12</b>	<b>0.6111111111111111</b>
13	0.6111111111111111
14	0.6111111111111111
15	0.6111111111111111

We find the best performance at max\_depth=12. Also, the performance keeps to be 0.6111 even max\_depth increases. Thus, we can conclude that the fully developed decision tree has a max\_depth of 12.

### **k-Nearest Neighbor:**

Settings:

KNeighborsClassifier(n\_neighbors=1, weights= 'distance')

We set weights to be 'distance' instead of 'uniform' because we put the weights into consideration, which can provide a better performance.

To find the n\_neighbors:

n_neighbors	weights= 'distance'	weights= 'uniform'
1	<b>0.1744421906693712</b>	<b>0.1744421906693712</b>
2	0.1744421906693712	0.06039076376554174
3	0.14052697616060225	0.128686327077748
4	0.13559322033898305	0.07692307692307693
5	0.12446958981612447	0.12037037037037038
6	0.13037037037037036	0.06427221172022685
7	0.0947867298578199	0.09059233449477352

We find the n\_neighbors by ourselves because GridSearchCV seems to be inaccurate for k-nearest neighbors using training dataset as sample. We find that n\_neighbors=1 gives the best performance.

## SVM

SVC(C=6, kernel='linear', random\_state=0)

Finding the best C:

C	Performance (linear)
1	0.15177065767284992
2	0.15075376884422112
3	0.15126050420168066
4	0.15151515151515152
5	0.14310051107325383
6	<b>0.1602671118530885</b>
7	0.15228426395939085
8	0.15640599001663896
9	0.14839797639123103
10	0.14625850340136054
11	0.15488215488215487
15	0.14625850340136054
30	0.14864864864864866

We find that the best performance is at C=6.

**Bayes:**

Settings:

CategoricalNB(alpha=0, fit\_prior=True, class\_prior=None)

The reason for choosing CategoricalNB is mentioned in Task 1 Bayes.

Finding the best alpha: (set fit\_prior=True, class\_prior=None)

alpha	Performance
<b>0</b>	<b>0.5109374999999999</b>
1	0.507537688442211
2	0.4967658473479948
3	0.5006587615283268
4	0.4899057873485868
5	0.48707482993197276
6	0.484931506849315
7	0.4834254143646408
8	0.48543689320388345
9	0.4804469273743016
10	0.4797768479776847
15	0.4663805436337625
30	0.44052863436123346
50	0.39694656488549623
80	0.3653543307086614
100	0.33599999999999997

We find the alpha=0 gives the best performance.

Then, we look at the parameter 'fit\_prior'.

fit_prior=False	0.41870350690754515
<b>fit_prior=True</b>	<b>0.5109374999999999</b>

We find that alpha=0 and fit\_prior=True gives the best performance with a value of 0.5109374999999999.

alpha=0 means that Laplace is not implemented. fit\_prior=True means that the classes have prior probabilities rather than uniform prior. For class\_prior, we would like to modify the class prior probabilities according to the data, so we set it to None. Therefore, we finally have the setting as CategoricalNB(alpha=0, fit\_prior=True, class\_prior=None).

## Bagging

Settings:

```
D = tree.DecisionTreeClassifier(min_samples_split=39, random_state=0)
BaggingClassifier(base_estimator=D, n_estimators=10, random_state=4,
bootstrap=True)
```

We have used DecisionTree, KNN, Bayes and SVM as the base\_estimator. We find that Bayes performs the best.

Random state is set in order to eliminate the randomness of the performance. And we find that the performance is slightly improved when random\_state is set to 4. bootstrap=True means that the samples are drawn with replacement.

We have tried all the combinations between min\_samples\_split from 2 to 100 and n\_estimators from 1 to 100. And we finally find the best performance at min\_samples\_split=39 and n\_estimators=10, with a performance of 0.630071599045346.

## Additional Modification:

Apart from the classifier tools, we do modification on the input data.

We first drop the columns that are unimportant.

```
csv = csv.drop(columns=['Exited'])
csv = csv.drop(columns=['RowNumber'])
csv = csv.drop(columns=['CustomerId'])
csv = csv.drop(columns=['Surname'])
```

Then, we group the data using 'pd.DataFrame'.

```
csv = pd.DataFrame(pd.get_dummies(csv))
```

All the features with non-numeric values will be changed to numeric, for example, for gender, male is changed to 0 while female is changed to 1. Besides, some features are separated out to form a new group. For instance, Geography with 3 choices (France, Germany, Spain) will be converted to three separate groups: 'France', 'Germany', 'Spain'. Customers from 'France' will have value 1 in 'France', value 0 in 'Germany' and value 0 in 'Spain'.

Lastly, we have observed some relationships and modified the 'Balance'.

```
for i in range(len(csv['Balance'])):
    if((csv['Balance'][i] == 0) and (csv['Tenure'][i] > 8)):
        csv['Balance'][i] = csv['Balance'].mean() +
csv['Balance'].std()*1
    if((csv['Balance'][i] == 0) and (csv['Tenure'][i] > 5)):
        csv['Balance'][i] = csv['Balance'].mean()

csv['Balance'] = preprocessing.minmax_scale(
    csv[['Balance']], feature_range=(0, 1))
```

When 'Balance'=0 and 'Tenure'>8, 'Balance' is changed to the mean+1sd of 'Balance'.

When 'Balance'=0 and 'Tenure'>5, 'Balance' is changed to the mean of 'Balance'.

We find the relationships by looking at Balance, Tenure and estimated\_salary. In normal sense, Balance should not be 0 when Tenure is relatively long, which means the customers are expected to have a higher Balance. Thus, we try to increase the Balance. We have found that the longer the Tenure, the higher the estimated salaries. Therefore, we modify Balance to mean+1sd for 'Tenure>8' and mean for '8>Tenure>5' respectively. For the conditions, we have tried different values and finally come out with 8 and 5.